

## More Register Allocation

### Last time

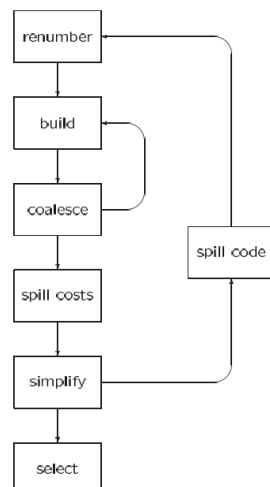
- Register allocation
  - Global allocation via graph coloring

### Today

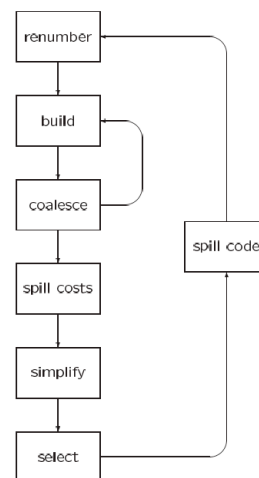
- More register allocation
  - Clarifications from last time
  - Finish improvements on basic graph coloring concept
  - Procedure calls
  - Interprocedural

## Interference Graph Allocators

### Chaitin



### Briggs



## Coalescing

### Move instructions


- Code generation can produce unnecessary move instructions  
`mov t1, t2`
- If we can assign `t1` and `t2` to the same register, we can eliminate the move

### Idea

- If `t1` and `t2` are not connected in the interference graph, **coalesce** them into a single variable

### Problem

- Coalescing can increase the number of edges and make a graph uncolorable
- Limit coalescing to avoid uncolorable graphs



CS553 Lecture


Register Allocation II

4

## Coalescing Logistics

### Rule

- If the virtual registers `s1` and `s2` do not interfere and there is a copy statement `s1 = s2` then `s1` and `s2` can be coalesced
- Steps
  - SSA
  - Find webs
  - Virtual registers
  - Interference graph
  - Coalesce



CS553 Lecture


Register Allocation II

5

### Example (Apply Chaitin algorithm)

Attempt to 3-color this graph (  ,  ,  )

Stack:  
d  
c  
b  
a<sub>2</sub>  
a<sub>1</sub>  
e



Weighted order:  
e  
a<sub>1</sub>  
a<sub>2</sub>  
b  
c  
d


The example indicates that nodes are visited in increasing weight order.  
Chaitin and Briggs visit nodes in an arbitrary order.

### Example (Apply Briggs Algorithm)

Attempt to 2-color this graph (  ,  )

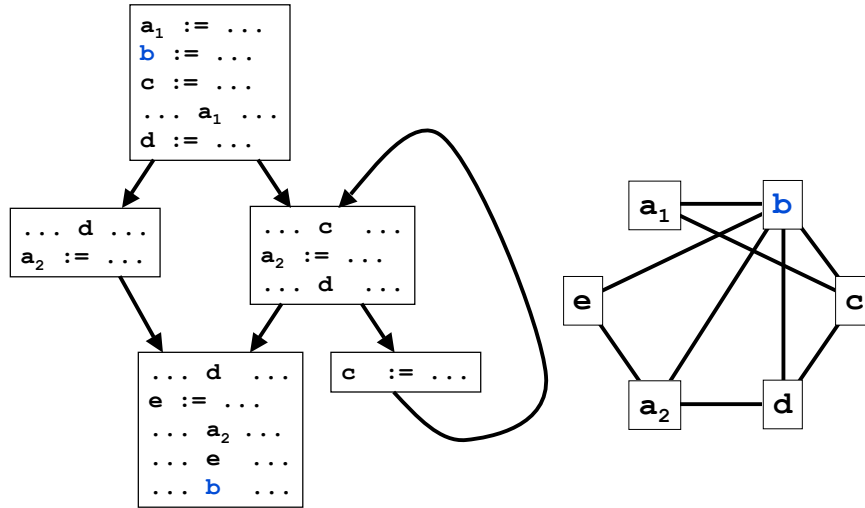
Stack:  
d  
c  
b\*  
a<sub>2</sub>\*  
a<sub>1</sub>\*  
e\*

\* blocked node



Weighted order:  
e  
a<sub>1</sub>  
a<sub>2</sub>  
b  
c  
d

## Spilling (Original CFG and Interference Graph)

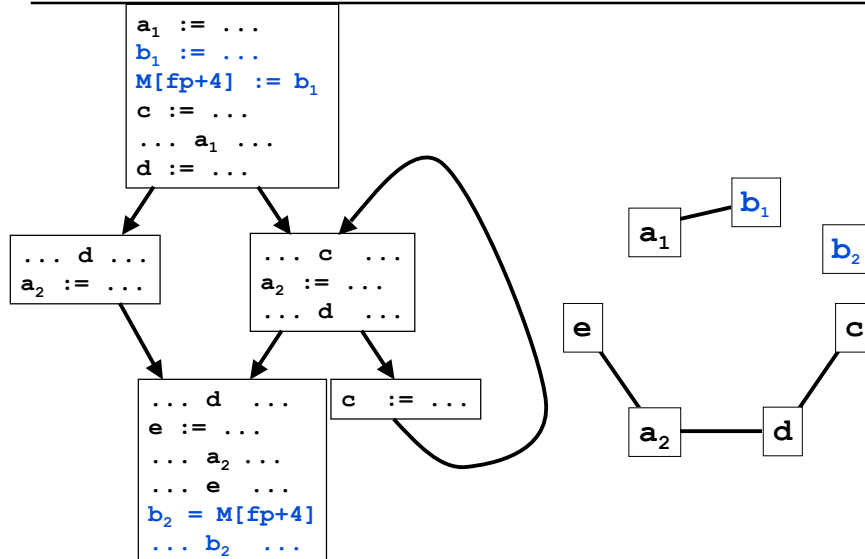


CS553 Lecture

Register Allocation II

8

## Spilling (After spilling b)



CS553 Lecture

Register Allocation II

9

### **Improvement #3: Live Range Splitting [Chow & Hennessy 84]**

---

#### **Idea**

- Start with variables as our allocation unit
- When a variable can't be allocated, split it into multiple subranges for separate allocation
- Selective spilling: put some subranges in registers, some in memory
- Insert memory operations at boundaries

#### **Why is this a good idea?**

### **Improvement #4: Rematerialization [Chaitin 82]&[Briggs 84]**

---

#### **Idea**

- Selectively re-compute values rather than loading from memory
- “Reverse CSE”

#### **Easy case**

- Value can be computed in single instruction, and
- All operands are available

#### **Examples**

- Constants
- Addresses of global variables
- Addresses of local variables (on stack)

## Complexity of Global Register Allocators

---

### Fastest to slowest

- Linear scan register allocation (Traub, Holloway, and Smith)
- Splitting allocators (Chow and Hennessey)
- Interference Allocator (Chaitin)
- Interference Allocator (Briggs)

### Interference Allocators

- Interference Graph construction:  $O(n^2)$  where  $n$  is the number of live ranges or webs

## Register Allocation and Procedure Calls

---

### Problem

- Register values may change across procedure calls
- The allocator must be sensitive to this

### Two approaches

- Work within a well-defined calling convention
- Use interprocedural allocation

## Calling Conventions

### Goals

- Fast calls (pass arguments in registers, minimal register saving/restoring)
- Language-independent
- Support debugging, profiling, *etc.*

### Complicating Issues

- Varargs
- Passing/returning aggregates
- Exceptions, non-local returns
  - `setjmp()`/`longjmp()`

## Architecture Review: Caller- and Callee-Saved Registers


### Partition registers into two categories

- Caller-saved
- Callee-saved

### Caller-saved registers

- Caller must save/restore these registers when live across call
- Callee is free to use them

### Example

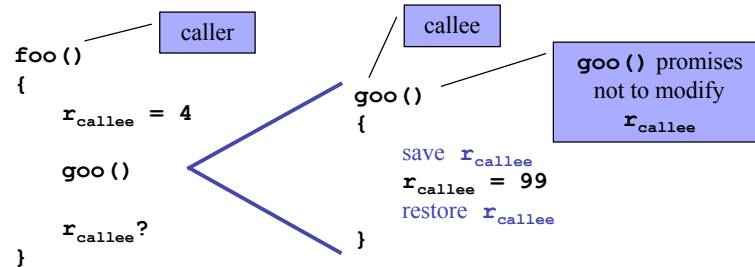


## Architecture Review: Caller- and Callee-Saved Registers

### Callee-saved registers

- Callee must save/restore these registers when it uses them
- Caller expects callee to not change them

### Example



## Register Allocation and Calling Conventions

### Insensitive register allocation

- Save all live caller-saved registers before call; restore after
- Save all used callee-saved registers at procedure entry; restore at return
- Suboptimal

```
foo()
{
    t = ...
    ... = t
    s = ...
    f()
    g()
    ... = s
}
```

A variable that is not live across calls should go in caller-saved registers

A variable that is live across multiple calls should go in callee-saved registers

### Sensitive register allocation

- Encode calling convention constraints in the IR and interference graph
- How? Use precolored nodes




## Precolored Nodes

### Add architectural registers to interference graph

- Precolored (mutually interfering)
- Not simplifiable
- Not spillable (infinite degree)

### Express allocation constraints

- Integers usually can't be stored in floating point registers
- Some instructions can only store result in certain registers
- Caller-saved and callee-saved registers. . .



CS553 Lecture

Register Allocation II

18

## Precolored Nodes and Calling Conventions

### Callee-saved registers

- Treat entry as def of all callee-saved registers
- Treat exit as use of them all
- Allocator must “spill” callee-saved registers to use them



### Caller-saved registers

- Variables live across call interfere with all caller-saved registers
- Splitting can be used (before/during/after call segments)

CS553 Lecture


Register Allocation II

19

## Example

```
foo() :  
  def(r3)  
  t1 := r3  
  a := ...  
  b := ...  
  ... a ...  
  call goo  
  ... b ...  
  r3 := t1  
  use(r3)  
  return
```

r1, r2 caller-saved  
r3 callee-saved



## Tradeoffs

### Callee-saved registers

- + Decreases code size: one procedure body may have multiple calls
- + Small procedures tend to need fewer registers than large ones; callee-save makes sense because procedure sizes are shrinking
- May increase execution time: For long-lived variables, may save and restore registers multiple times, once for each procedure, instead of a single end-to-end save/restore

### The larger “problem”

- We’re making local decisions for policies that require global information

## Interprocedural Register Allocation

---

### Wouldn't it be nice to . . .

- Allocate registers across calls to minimize unnecessary saves/restores?
- Allocate global variables to registers over entire program?

### Compile-time interprocedural register allocation?

- + Could have great performance
- Might be expensive
- Might require lots of recompilation after changes (no separate compilation?)

### Link-time interprocedural re-allocation?

- + Low compile-time cost
- + Little impact on separate compilation
- Link-time cost

## Wall's Link-time Register Allocator [Wall 86]

---

### Overall strategy

- Compiler uses 8 registers for local register allocation
- Linker controls allocation of remaining 52 registers

### Compiler does local allocation & planning for linker

- Load all values at beginning of each basic block; store all values at end of each basic block
- Generate call graph information
- Generate variable usage information for each procedure
- Generate **register actions**

### Linker does interprocedural allocation & patches compiled code

- Generates “interference graph” among variables
- Picks best variables to allocate to registers
- Executes register actions for allocated variables to patch code

## Register Actions

---

Describe code patch if particular variable allocated to a register

- **REMOVE(var)**: Delete instruction if **var** allocated to a register
- **OPx(var)**: Replace op x with register that was allocated to **var**
- **RESULT(var)**: Replace result with register allocated to **var**

Usage

```
-r := load var: REMOVE(var)

-ri := rj op rk:
    OP1(var)      if var loaded into rj
    OP2(var)      if var loaded into rk
    RESULT(var)   if var stored from ri

-store var := r: REMOVE(var)
```

## Example

---

```
w := (x + y) * z
```

r1 := load x	REMOVE(x)
r2 := load y	REMOVE(y)
r3 := r1 + r2	OP1(x), OP2(y)
r4 := load z	REMOVE(z)
r5 := r3 * r4	OP2(z), RESULT(w)
store w := r5	REMOVE(w)

## Another Example

**w := y++ \* z**

Suppose **y** is allocated to register **r5**

	<b>REMOVE(y)</b>
<b>r5 := r5 + 1</b>	<b>OP1(y), RESULT(y)</b>
	<b>REMOVE(y)</b>
<b>r2 := load z</b>	<b>REMOVE(z)</b>
<b>r1 := r5 * r2</b>	<b>OP1(y), OP2(z), RESULT(w)</b>
<b>store w := r1</b>	<b>REMOVE(w)</b>

### Problem

- Loaded value is still live after store overwrites it
- Post-incremented value of **y** is lost if **y** is allocated to register
- We need to registers to hold the two values of **y**

## Extension

### More actions

- **LOAD(var)**: Replace load with move from the register holding **var**
- **STORE(var)**: Replace store with move to the register holding **var**

### **LOAD(var)**

- Use instead of **REMOVE(var)** if **var** is stored into while result of load is still live

### **STORE(var)**

- Use instead of **REMOVE(var)** if source is stored into more than one variable

## Example Revisited


**w := x := y++ \* z**

<b>r1 := load y</b>	REMOVE(y)	<b>LOAD(y)</b>
<b>r2 := r1 + 1</b>	OP1(y), RESULT(y)	<b>RESULT(y)</b>
<b>store y := r2</b>	REMOVE(y)	<b>REMOVE(y)</b>
<b>r2 := load z</b>	REMOVE(z)	<b>REMOVE(z)</b>
<b>r1 := r1 * r2</b>	OP1(y), OP2(z), RESULT(w)	<b>OP2(z), RESULT(w)</b>
<b>store x := r1</b>		<b>STORE(x), OP1(w)</b>
<b>store w := r1</b>	REMOVE(w)	<b>REMOVE(w)</b>

## Deciding Which Variables to Promote to Registers

### Steps

- Use bottom-up algorithm to assign pseudo registers
- Allocate pseudo registers to non-simultaneously live variables
- Allocate real registers to most frequently used pseudo registers



## Possible Improvements

---

- Use profile data to construct weights**
- Do global register allocation at compile-time**
- Track liveness information for variables at each call site**
- Track intraprocedural interference graph**
- Use real interference graph at link-time**

## Performance Summary

---

**Machine: DEC WRL Titan RISC processor (64 registers)**

**Basic experiment**

- Local compile-time allocator uses 8 registers
  - Link-time allocator uses 52 registers
  - Simple static frequency estimates
  - Small benchmarks
- ⇒ 10-25% speed-up over local allocation alone

**Improvements**

- 0-6% with profile data
- 0-5% with compile-time global allocation

**Benefit decreases with number of link-time registers**

**Link-time better than global register allocation**

## Link-Time Register Allocation: The Big Picture


### Delayed decision making

- We can often make more informed decisions later in the translation process (link time)
- Requires communication among different system components, in this case the compiler and the linker
- Intuitively, more information is better, but effectively using this information can require cleverness

## JIT Environment

### Dynamic compilation requires fast register allocation

- Linear Scan Register Allocation [Poletto & Sarkar99]
- Not based on graph coloring
- Greedy algorithm based on live intervals
  - Spill the variable whose interval ends furthest in the future



- What if we had spilled a or b instead of c?



## Linear Scan Register Allocation

---

### Performance results

- Linear scan is linear in number of variables
- Graph coloring is  $O(n^2)$
- Code quality is within 12% of graph coloring

## Register Allocation Summary

---

### Simple solutions can work well

- Callee-save and caller-save registers
- Interference graph can be used to represent these constraints

### Delayed decision making

- Link-time register allocation is an instance of a theme:
  - Make decisions when more information is available
  - Staged compilation

### Modern environments change the rules

- Java is often dynamically compiled
  - Compilation time is important
  - Need to find fast and effective solutions

## Concepts

---

### Register allocation and procedure calls

#### Calling conventions

- Caller- vs. callee-saved registers
- Precoloring
- Finding register values in stack can be hard

#### Interprocedural analysis

- Link-time register allocation
  - Register actions

#### Register allocation in a JIT

- Linear Scan Register Allocation

## Next Time

---

### Project 2

- Due Monday at the beginning of class

### Lecture

- Code scheduling