

## Profile-Guided Optimizations

---

### Last time

- Instruction scheduling
  - Register renaming
  - Balanced Load Scheduling
  - Loop unrolling
  - Software pipelining

### Today

- More instruction scheduling
  - Profiling
  - Trace scheduling

## Motivation for Profiling


---

### Limitations of static analysis

- Compilers can analyze possible paths but must behave conservatively
- Frequency information cannot be obtained through static analysis

### How runtime information helps

- Control flow information



Optimize the more frequent path  
(perhaps at the expense of the less frequent path)


- Memory conflicts

```
st r1, 0(r5) } If r5 and r4 always have different values,  
ld r2, 0(r4) } we can move the load above the store
```

## Profile-Guided Optimizations

---

### Basic idea

- Instrument and run program on sample inputs to get likely runtime behavior
-  – Can use this information to improve instruction scheduling
- Many other uses
  - Code placement
  - Inlining
  - Value speculation
  - Branch prediction
  - Class-based optimization (static method lookup)

## Profiling Issues

---


### Profile data

- Collected over whole program run
- May not be useful (unbiased branches)
- May not reflect all runs
- May be expensive and inconvenient to gather
  - Continuous profiling [Anderson 97]
- May interfere with program

## Control-Flow Profiles

Commonly gather two types of information

- **Execution frequencies** of basic blocks
- **Branch frequencies** of conditional branches
- Represent information in a **weighted flow graph**




**Instrumentation**

- Insert instrumentation code at basic block entrances and before each branch
- Take average of values from multiple training runs
- Fairly inexpensive


## Code Motion Using Control Flow Profiles

**Code motion across basic blocks**

- Increased scheduling freedom



- If we want to move s1 to A, we must move s1 to both A and B






- If we want to move s1 to B, we must move s1 to both B and C

## Code Motion Using Control Flow Profiles (cont)

### Code motion across basic blocks

- Increased scheduling freedom



- If we want to move s1 from B to A and if s1 would destroy a value along the A→C path, do renaming (in hardware or software)
- What if s1 might cause an exception?

## Memory-Dependence Profiles

### Gather information about memory conflicts

- Frequencies of **address matches** between pairs of loads and stores
- Attempts to answer the question: Are two references independent of one another?
- Concentrate on **ambiguous** reference pairs (those that the compiler cannot figure out)

st1: store r5  
ld2: load r4

(st1, ld2, 7) If this number is low, we can speculatively assume that st1 and ld2 do not conflict

### Instrumentation

- Much more expensive (in both space and time) to gather than control flow information
- First perform control flow profiling
- Apply only to most frequently executed blocks

## Trace Scheduling [Fisher 81] and [Ellis 85]

### Basic idea


- We want large blocks to create large scheduling windows, but basic blocks are small because branches are frequent
- Create **superblocks** to increase scheduling window
- Use profile information to create good superblocks
- Optimize each superblock independently

### Superblocks

- A sequence of basic blocks with a single entrance and multiple exits

### Goals

- Want large superblocks
- Want to avoid early exits
- Want blocks that match actual execution paths



## Trace Scheduling (example)

```
b[i] = "old"
a[i] = ...
if (a[i]>0) then
    b[i]="new";
else
    stmt X
    stmt Y
endif
c[i] = ...
```

```
trace:
b[i] = "old"
a[i] = ...
b[i]="new";
c[i] = ...
if (a[i]<=0) then goto repair
continue:
...
```

```
repair:
restore old b[i]
stmt X
stmt Y
recalculate c[i]?
goto continue
```


## Trace Scheduling (cont)

### Three steps

1. Create superblocks
2. Enlarge superblocks
3. Compact (optimize) superblocks

### 1. Superblock formation

- Create **traces** using **mutual-most-likely** heuristic (two blocks **A** and **B** are mutual-most-likely if **B** is the most likely successor of **A**, and **A** is the most likely predecessor of **B**)




– A **trace** is a maximal sequence of mutual-most-likely blocks that does not contain a back edge

– Each block belongs to exactly one trace

## Trace Scheduling (cont)

### 1. Superblock formation (cont)

- Convert traces into Superblocks
- Use tail duplication to eliminate side entrances



- Tail duplication increases code size

## Trace Scheduling (cont)

---

### 2. Superblock enlargement

- Enlarge superblocks that are too small
- Code expansion can hurt i-cache performance

### Three techniques for enlargement

- **Branch target expansion**
  - If the last branch in a superblock is likely to jump to the start of another superblock, append the contents of the target superblock to the first superblock
- **Loop peeling**
- **Loop unrolling**
  - These last two techniques apply to **superblock loops**, which are superblocks whose last blocks are likely to jump to their first blocks
  - Assume that each loop body has a single dominant path

## Trace Scheduling (cont)

---

### 3. Optimizations

- Perform list scheduling for each superblock
- Memory-dependence profiles can be used to speculatively assume that load/store pairs do not conflict
  - Insert repair code in case the assumption is incorrect
- Software pipelining

### Speculation based on memory-dependence profiles (example)

```
b[i] = "old"
a[i] = ...
if (a[i]>0) then
  b[i]="new";
else
  stmt X
  stmt Y
endif
c[i] = a[j]
```

```
trace:
b[i] = "old"
c[i] = a[j]
a[i] = ...
b[i]="new";
if (i==j) then goto deprepair
if (a[i]<=0) then goto repair
continue:
...
```

```
deprepair:
c[i] = a[i]
if (a[i]<=0) then goto repair
goto continue
```

```
repair:
restore old b[i]
stmt X
stmt Y
goto continue
```

CS553 Lecture


Profile-Guided Optimizations

16


### Enhancements to Profile-Guided Code Scheduling

#### Path profiling [Ball and Larus 96]


- Collect information about entire paths instead of about individual edges



Edge profiles



Path profiles



Path profiles

- Limit paths to some specified length (can thus handle loops)
- Can also stop paths at back edges
- Disadvantages of path profiling?

CS553 Lecture

Profile-Guided Optimizations

17



## Lessons

---

### **Larger scope helps**

- How can we increase scope? How do we schedule across control dependences?

### **Static information is limited**

- Use profiles
- How else can profiles be used in optimization?
- Can we do these kinds of optimizations at runtime?

## Concepts

---

### **Instruction scheduling**

- Software pipelining
- Trace scheduling
- Both use profile information
- Both look at scopes beyond basic blocks

### **Miscellany**

- Path profiling
- Speculative hedging

## Next Time

---

### Reading

- Mahlke'92

### Lecture

- Speculation and predication