

CS 243

Lecture 11

Binary Decision Diagrams (BDDs)

in Pointer Analysis

1. Datalog \rightarrow BDD
2. BDDs
3. Context-Sensitive Pointer Analysis
4. Performance of BDD Algorithms

Readings: Chapter 12

Automatic Analysis Generation



Programmer:
Security analysis
in 10 lines

Compiler writer:
Ptr analysis
in 10 lines

1000s of lines
1 year tuning

PQL

Datalog


bddbddb
(**BDD-based**
deductive database)
with
Active Machine Learning

BDD operations

BDD: 10,000s-lines library

1. Datalog \rightarrow BDDs

- Example




calls(A,B)
calls(A,C)
calls(A,D)
calls(B,D)
calls(C,D)

Call Graph Relation


x_1	x_2	x_3	x_4	f
0	0	0	0	0
0	0	0	1	1
0	0	1	0	1
0	0	1	1	1
0	1	0	0	0
0	1	0	1	0
0	1	1	0	0
0	1	1	1	1
1	0	0	0	0
1	0	0	1	0
1	0	1	0	0
1	0	1	1	1
1	1	0	0	0
1	1	0	1	0
1	1	1	0	0
1	1	1	1	0

- Relation expressed as a binary function.
 - $A=00$, $B=01$, $C=10$, $D=11$




Binary Decision Diagrams (Bryant, 1986)

- Graphical encoding of a truth table.




Binary Decision Diagrams

- Collapse redundant nodes.




Binary Decision Diagrams

- Collapse redundant nodes.




Binary Decision Diagrams

- Collapse redundant nodes.




Binary Decision Diagrams

- Collapse redundant nodes.




Binary Decision Diagrams

- Eliminate unnecessary nodes.



Binary Decision Diagrams

- Eliminate unnecessary nodes.



Datalog \rightarrow BDDs


Datalog	BDDs
Relations	Boolean functions
Relation ops: \bowtie , \cup , select, project	Boolean function ops: \wedge , \vee , \neg , \sim
Relation at a time	Function at a time
Semi-naïve evaluation	Incrementalization
Fixed-point	Iterate until stable

2. Binary Decision Diagrams


- Represent tiny and huge relations compactly
- Size depends on redundancy
 - Similar contexts have similar numberings
 - Variable ordering in BDDs

BDD Variable Order is Important!

$$x_1x_2 + x_3x_4$$





$$x_1 < x_2 < x_3 < x_4$$





$$x_1 < x_3 < x_2 < x_4$$

Expanded Call Graph



Numbering Clones



3. Context-Sensitive Pointer Analysis Algorithm

1. First, do context-insensitive pointer analysis to get call graph.
2. Number clones.
3. Do context-insensitive algorithm on the cloned graph.
 - Results explicitly generated for every clone.
 - Individual results retrievable with Datalog query.

4. Performance of BDD Algorithm

- Direct implementation
 - Does not finish even for small programs
 - > 3000 lines of code
- Requires tuning for about 1 year
- Easy to make mistakes
 - Mistakes found months later

An Adventure in BDDs

- Context-sensitive numbering scheme
 - Modify BDD library to add special operations.
 - Can't even analyze small programs. *Time: ∞*
- Improved variable ordering
 - Group similar BDD variables together.
 - Interleave equivalence relations.
 - Move common subsets to edges of variable order. *Time: 40h*
- Incrementalize outermost loop
 - Very tricky, many bugs. *Time: 36h*
- Factor away control flow, assignments
 - Reduces number of variables *Time: 32h*

An Adventure in BDDs

- Exhaustive search for best BDD order
 - Limit search space by not considering intradomain orderings. *Time: 10h*
- Eliminate expensive rename operations
 - When rename changes relative order, result is not isomorphic. *Time: 7h*
- Improved BDD memory layout
 - Preallocate to guarantee contiguous. *Time: 6h*
- BDD operation cache tuning
 - Too small: redo work, too big: bad locality
 - Parameter sweep to find best values. *Time: 2h*

An Adventure in BDDs

- Simplified treatment of exceptions
 - Reduce number of vars, iterations necessary for convergence. *Time: 1h*
- Change iteration order
 - Required redoing much of the code. *Time: 48m*
- Eliminate redundant operations
 - Introduced subtle bugs. *Time: 45m*
- Specialized caches for different operations
 - Different caches for and, or, etc. *Time: 41m*

An Adventure in BDDs

- Compacted BDD nodes
 - 20 bytes → 16 bytes *Time: 38m*
- Improved BDD hashing function
 - Simpler hash function. *Time: 37m*
- Total development time: 1 year
 - 1 year per analysis?!?
- Optimizations obscured the algorithm.
- Many bugs discovered, maybe still more.

Variable Numbering: Active Machine Learning

- Must be determined dynamically
- Limit trials with properties of relations
- Each trial may take a long time
- Active learning:
select trials based on uncertainty
- Several hours
- Comparable to exhaustive for small apps

Optimizations in bddbddb

- Algorithmic
 - Clever context numbering to exploit similarities
- Query optimizations
 - Magic-set transformation
 - semi-naïve evaluation
- Compiler optimizations
 - Redundancy elimination, liveness analysis
- BDD optimizations
 - Active machine learning
- BDD library extensions and turning

Automatic Analysis Generation



Programmer:
Security analysis
in 10 lines

Compiler writer:
Ptr analysis
in 10 lines

1000s of lines
1 year tuning

PQL

Datalog

BDD operations

bddbddb
(**BDD**-based
deductive database)
with
Active Machine Learning

BDD: 10,000s-lines library

Software

- System is publicly available at:
<http://bddbddb.sourceforge.net>