

Session: Systems for ML 1

Moderator: **Martin Maas**

Talk Order:

Oral: MLPerf Mobile Inference Benchmark: An Industry-Standard Open-Source Machine Learning Inference Benchmark

Oral: TorchSparse: Efficient Point Cloud Inference Engine

Oral: Hydrozoa: Dynamic Hybrid-Parallel DNN Training on Serverless Containers

Oral: Plumber: Diagnosing and Removing Performance Bottlenecks in Machine Learning Data Pipelines

Oral: Graphiler: Optimizing Graph Neural Networks with Message Passing Data Flow Graph

View on <https://mlsys.org>



MLPerf Mobile Inference Benchmark: An Industry-Standard Open-Source Machine Learning Benchmark for On-Device AI

- Vijay Janapa Reddi
- David Kanter
- Peter Mattson
- Jared Duke
- Thai Nguyen
- Ramesh Chukka
- Ken Shiring
- Koan-Sin Tan
- Mark Charlebois
- William Chou
- Mostafa El-Khamy

- Jungwook Hong
 - Tom St John
 - Cindy Trinh
 - Michael Buch
 - Mark Mazumder
 - Relja Markovic
 - Thomas Atta
 - Fatih Cakir
 - Masoud Charkhabi
 - Xiaodong Chen
 - Cheng-Ming Chiang
 - Dave Dexter
 - Terry Heo
 - Guenther Schmuelling
 - Maryam Shabani
 - Dylan Zika
- MLSys 2022

View on <https://mlsys.org>



TorchSparse: Efficient Point Cloud Inference Engine

- Haotian Tang
- Zhijian Liu
- Xiuyu Li
- Yujun Lin
- Song Han

View on <https://mlsys.org>



MLSys 2022

Hydrozoa: Dynamic Hybrid-Parallel DNN Training on Serverless Containers

- Runsheng Guo
- Victor Guo
- Antonio Kim
- Josh Hildred
- Khuzaima Daudjee

View on <https://mlsys.org>



Plumber: Diagnosing and Removing Performance Bottlenecks in Machine Learning Data Pipelines

- Michael Kuchnik
- Ana Klimovic
- Jiri Simsa
- Virginia Smith
- George Amvrosiadis

View on <https://mlsys.org>



Graphiler: Optimizing Graph Neural Networks with Message Passing Data Flow Graph

- Zhiqiang Xie
- Minjie Wang
- Zihao Ye
- Zheng Zhang
- Rui Fan

View on <https://mlsys.org>

